

KNOWLEDGE CHECK

Databricks Data Engineer Associate

Course quizzes for all three days, with answer key

The quiz has 68 multiple-choice questions covering the three course days: 20 for Day 1, 20 for Day 2 and 28 for Day 3. Questions are numbered continuously and follow the style of the certification exam. Choose the single best answer to each question. The answer key is on the last page.

DAY 1

Day 1 – Platform, ingestion and Delta Lake

1. Which component of the Databricks Lakehouse Platform serves as the centralized governance layer for data, users, and permissions?
 - A. Databricks Runtime
 - B. Delta Lake
 - C. Unity Catalog
 - D. Photon Engine
2. In Unity Catalog, what is the correct hierarchy of securable objects from broadest to most granular?
 - A. Catalog > Schema > Table
 - B. Schema > Catalog > Table
 - C. Metastore > Schema > Catalog
 - D. Workspace > Catalog > Table
3. A user who does not own any objects must query the table `sales.bronze.orders` in Unity Catalog. What is the minimum set of privileges they need?
 - A. Only `SELECT` on the table
 - B. `USE CATALOG` on `sales` and `USE SCHEMA` on `sales.bronze` only
 - C. `USE CATALOG` on `sales`, `USE SCHEMA` on `sales.bronze` and `SELECT` on the table
 - D. `BROWSE` on the catalog
4. Which SQL command reads data from a CSV file in a Unity Catalog Volume and creates a managed Delta table?
 - A. `LOAD DATA INPATH '/path' INTO TABLE t`
 - B. `COPY INTO t FROM '/path'`
 - C. `INSERT INTO t VALUES (SELECT * FROM '/path')`
 - D. `CREATE TABLE t AS SELECT * FROM read_files('/path', format => 'csv')`
5. Which statement about managed vs external tables in Unity Catalog is correct?
 - A. Managed tables store data in a user-specified cloud location
 - B. External tables have their data deleted when the table is dropped
 - C. Managed tables store data in the catalog's managed storage and data is deleted on `DROP`
 - D. External tables cannot be registered in Unity Catalog

6. A data engineer wants to apply SQL transformations on a DataFrame `df`. Which is the correct sequence?

A.

```
df.createOrReplaceTempView("v")
spark.sql("SELECT * FROM v")
```

B.

```
df.registerTable("v")
spark.sql("SELECT * FROM v")
```

C.

```
df.saveAsTable("v")
spark.sql("SELECT * FROM v")
```

D. `spark.sql("CREATE VIEW v AS df")`

7. What does Delta Lake's transaction log guarantee?

- A. Exactly-once semantics for batch reads only
- B. ACID transactions, time travel, and schema enforcement
- C. Automatic data compression only
- D. Distributed caching of query results

8. Which file format does Delta Lake use to store data?

- A. Parquet
- B. ORC
- C. JSON
- D. Avro

9. A `DESCRIBE HISTORY` command on a Delta table shows 5 versions. What does each version represent?

- A. A schema change only
- B. A transactional operation (write, update, delete, merge, optimize, etc.)
- C. A user session
- D. A scheduled job execution

10. A data engineer runs the following query. What does it return?

```
SELECT * FROM my_table VERSION AS OF 3
```

- A. The table as it existed at transaction version 3
- B. The current table filtered to 3 rows
- C. The 3rd partition of the table
- D. An error because time travel is not supported

11. A team of 40 business analysts runs SQL queries and Power BI dashboards concurrently during working hours, and nothing runs at night. They want fast start-up, automatic scaling for concurrency and no infrastructure to manage. Which compute is the most suitable?

- A. A GPU-enabled all-purpose cluster
- B. A single-node all-purpose cluster shared by all analysts
- C. A new job cluster created for every dashboard refresh
- D. A serverless SQL warehouse

12. Which of the following is the recommended way to store non-tabular files (CSV, JSON, images) in Unity Catalog?

- A. DBFS root
- B. DBFS mounts
- C. Unity Catalog Volumes
- D. Cluster local storage

13. What is the difference between ELT and ETL in the context of Databricks Lakehouse?

- A. They are the same process
- B. ELT is for batch only; ETL is for streaming
- C. ELT requires external tools; ETL is native
- D. ELT loads raw data first then transforms in-place; ETL transforms data before loading

14. What does `DESCRIBE DETAIL my_table` return that `DESCRIBE TABLE my_table` does not?

- A. Column names and data types
- B. Table location, file count, size in bytes, partitioning info, and table properties
- C. The SQL definition of the table
- D. Access control permissions

15. On a classic cluster, a data engineer runs the following commands. What is the impact on time travel?

```
SET spark.databricks.delta.retentionDurationCheck.enabled = false;
VACUUM my_table RETAIN 0 HOURS;
```

- A. No impact — time travel still works for all versions
- B. An error — VACUUM requires at least 168 hours retention
- C. Only the latest version is vacuumed
- D. All historical data files are removed; time travel to old versions will fail

16. What does the `%run` magic command do in a Databricks notebook?

- A. Runs an external Python script on the driver node
- B. Executes another notebook in the same execution context, sharing variables and functions
- C. Runs a SQL query and returns results
- D. Restarts the cluster and runs all cells

17. The incoming DataFrame `df` has a column that does not exist in the Delta table `my_table`. What happens when this code runs?

```
df.write.mode("append").saveAsTable("my_table")
```

- A. The new column is silently ignored
- B. The entire table is recreated with the new schema
- C. The new column is automatically added to the table
- D. The write fails with a schema mismatch error (schema enforcement)

18. Which `dbutils` submodule provides methods for listing, copying, and removing files?

- A. `dbutils.fs.ls("/path")`
- B. `dbutils.data.list("/path")`
- C. `dbutils.files.browse("/path")`
- D. `dbutils.storage.ls("/path")`

19. What does `RESTORE TABLE my_table TO VERSION AS OF 2` do?

- A. Creates a copy of version 2 as a new table
- B. Deletes all versions after version 2
- C. Reverts the table to its state at version 2, creating a new version in the history
- D. Displays the table at version 2 without modifying it

20. Which statement correctly reads a JSON file with an explicit schema in PySpark?

- A. `df = spark.read.json("/path/data.json").schema("id INT, name STRING")`
- B.

```
from pyspark.sql.types import StructType, StructField, IntegerType, StringType
schema = StructType([StructField("id", IntegerType()), StructField("name", StringType())])
df = spark.read.schema(schema).json("/path/data.json")
```

- C. `df = spark.read.format("json").option("schema", "id:int, name:string").load("/path")`
- D. `df = spark.read.json("/path").cast("id INT, name STRING")`

DAY 2

Day 2 – Optimization, incremental ingestion and transformations

21. What does the `OPTIMIZE` command do on a Delta table?

- A. Compacts small files into larger, optimally-sized files
- B. Deletes old versions from the transaction log
- C. Adds indexes to all columns
- D. Converts the table from Parquet to Delta format

22. A join between a 400 MB `orders` table and a 25 MB `products` dimension runs as a sort-merge join with a large shuffle. The data engineer wants Spark to broadcast `products` automatically in this session and to confirm that the change actually helps. What should they do?

- A. Set `spark.sql.autoBroadcastJoinThreshold` to -1 and rerun the query
- B. Increase `spark.sql.shuffle.partitions` to 2000 and assume the join is now faster
- C. Raise `spark.sql.autoBroadcastJoinThreshold` above the size of `products` (for example 50 MB), rerun the query, and compare the physical plan and run time with the previous run
- D. Call `orders_df.cache()` before the join; no measurement is needed

23. What does `VACUUM` do on a Delta table?

- A. Removes duplicate rows
- B. Removes data files no longer referenced by the transaction log (older than retention threshold)
- C. Compresses the transaction log
- D. Recalculates table statistics

24. Which Delta Lake feature replaces traditional partitioning and Z-Ordering with automatic data layout optimization?

- A. Auto Optimize
- B. Predictive Optimization
- C. Liquid Clustering
- D. Adaptive Query Execution

25. An Auto Loader stream ingests JSON files from a landing location that now holds millions of files, with thousands of new files arriving every hour. Most of each micro-batch is spent discovering which files are new. Which change addresses this?

- A. Switch from directory listing mode to file notification mode, preferably with file events (`cloudFiles.useManagedFileEvents`), so new files are discovered from storage events instead of by listing the directory
- B. Replace Auto Loader with a scheduled `COPY INTO` that scans the same location
- C. Remove `cloudFiles.schemaLocation` so the schema is no longer tracked
- D. Set `cloudFiles.inferColumnTypes` to `false`

26. What is the difference between a Temporary View and a Permanent View in Databricks?

- A. Temporary views are stored in Unity Catalog; permanent views are not
- B. Temporary views exist only within the SparkSession; permanent views are persisted in the metastore
- C. Temporary views support SQL only; permanent views support PySpark only
- D. There is no difference

27. What is the role of a checkpoint directory in Structured Streaming?

- A. It stores the output data
- B. It caches the stream schema
- C. It tracks which data has already been processed to enable exactly-once semantics
- D. It logs user queries

28. Which Auto Loader option specifies the file format to ingest?

- A. `cloudFiles.path`
- B. `cloudFiles.format`
- C. `cloudFiles.schemaLocation`
- D. `cloudFiles.inferColumnTypes`

29. What does `trigger(availableNow=True)` do in a streaming query?

- A. Runs the stream continuously with micro-batches
- B. Runs exactly one micro-batch then pauses
- C. Waits for new data indefinitely
- D. Processes all available data incrementally and then stops

30. What is the difference between `COPY INTO` and Auto Loader for incremental file ingestion?

- A. `COPY INTO` uses file notification; Auto Loader uses directory listing
- B. `COPY INTO` stores state in a checkpoint; Auto Loader stores state in the Delta log
- C. `COPY INTO` tracks processed files in the Delta log (idempotent SQL); Auto Loader uses checkpoints and scales to millions of files
- D. There is no difference; they are aliases

31. Which of the following correctly shows an Auto Loader read stream?

A.

```
spark.readStream.format("cloudFiles") \
  .option("cloudFiles.format", "json") \
  .option("cloudFiles.schemaLocation", checkpoint) \
  .load("/source/path")
```

B.

```
spark.read.format("autoLoader") \
  .option("format", "json") \
  .load("/source/path")
```

C.

```
spark.readStream.format("autoLoader") \
  .option("fileFormat", "json") \
  .load("/source/path")
```

D.

```
spark.readStream.format("delta") \
  .option("autoLoader", True) \
  .load("/source/path")
```

32. What is the purpose of Predictive Optimization in Databricks?

- A. Predicts query execution time
- B. Predicts storage costs for the next month
- C. Optimizes cluster autoscaling
- D. Automatically runs OPTIMIZE, VACUUM, and ANALYZE TABLE based on table usage patterns

33. What do Deletion Vectors improve in Delta Lake?

- A. INSERT performance
- B. DELETE, UPDATE, and MERGE performance by marking rows as deleted without rewriting files
- C. SELECT performance for aggregations
- D. VACUUM speed

34. Which SQL construct creates a running total (cumulative sum) of `amount` ordered by `date` within each category?

- A. `SUM(amount) GROUP BY category`
- B. `CUMSUM(amount) WITHIN GROUP (ORDER BY date)`
- C. `RUNNING_SUM(amount, category, date)`
- D.

```
SUM(amount) OVER (  
  PARTITION BY category  
  ORDER BY date  
  ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW  
)
```

35. What does the `explode()` function do?

- A. Converts each element of an array or map into a separate row
- B. Splits a string by delimiter into multiple rows
- C. Flattens nested JSON into a flat schema
- D. Decompresses compressed columns

36. Two monthly extracts, `orders_jan` and `orders_feb`, have identical columns. A data engineer must combine them in SQL, keeping every row (including rows that appear in both), without paying for duplicate elimination. Which query should they use?

- A. `SELECT * FROM orders_jan UNION SELECT * FROM orders_feb`
- B. `SELECT * FROM orders_jan UNION ALL SELECT * FROM orders_feb`
- C. `SELECT * FROM orders_jan CROSS JOIN orders_feb`
- D. `SELECT * FROM orders_jan INNER JOIN orders_feb USING (order_id)`

37. Which window function assigns a unique sequential number to each row within a partition, with no gaps?

- A. `RANK()`
- B. `DENSE_RANK()`
- C. `NTILE()`
- D. `ROW_NUMBER()`

38. On a table with billions of rows, an analyst needs the number of distinct `customer_id` values per segment; a result within a few percent is acceptable if the query is fast. They also want count, mean, standard deviation, min, quartiles and max of `total_amount` from a single call. Which combination fits best?

- A. `count("customer_id")` and `df.select("total_amount").summary()`
- B. `count_distinct("customer_id")` and `df.select("total_amount").describe()`
- C. `approx_count_distinct("customer_id")` and `df.select("total_amount").summary()`
- D. `approx_count_distinct("customer_id")` and `df.select("total_amount").explain()`

39. In the Spark UI, a sort-merge join stage shows 199 tasks finishing in about 5 seconds and one task running for 14 minutes. In Summary Metrics, the maximum shuffle read is 9 GB while the median is 40 MB. What is the most likely cause, and which built-in mechanism is designed to handle it?

- A. Data skew on the join key; Adaptive Query Execution skew join handling splits the oversized partition into smaller tasks
- B. Too little memory on every executor; increase `spark.executor.memory` on all nodes
- C. Too many shuffle partitions; set `spark.sql.shuffle.partitions` to 1
- D. Too many small files in the source tables; run `VACUUM` on both tables

40. A data engineer has a DataFrame with duplicate rows. Which PySpark method removes duplicate rows based on a subset of columns?

- A. `df.distinct()`
- B. `df.unique(["customer_id"])`
- C. `df.removeDuplicates(["customer_id"])`
- D. `df.dropDuplicates(["customer_id", "email"])`

DAY 3

Day 3 – Pipelines, jobs, CI/CD, troubleshooting and governance

41. In the Medallion Architecture, which layer contains raw, unvalidated data ingested from source systems?

- A. Silver
- B. Gold
- C. Bronze
- D. Platinum

42. In a Lakeflow Spark Declarative Pipeline, what is the difference between a `STREAMING TABLE` and a `MATERIALIZED VIEW`?

- A. `STREAMING TABLE` supports batch only; `MATERIALIZED VIEW` supports streaming only
- B. `STREAMING TABLE` processes only the input rows that are new since the last update (append-only sources); `MATERIALIZED VIEW` keeps the stored result of its query up to date on each refresh
- C. They are identical; the names are interchangeable
- D. `STREAMING TABLE` stores views; `MATERIALIZED VIEW` stores tables

43. In Lakeflow Spark Declarative Pipelines (default publishing mode), how do you normally reference another dataset defined in the same pipeline?

- A. By the cloud storage path of the dataset's files
- B. With a `LIVE.` prefix, which is required in every pipeline
- C. With a `PIPELINE.` prefix
- D. By the dataset name directly (e.g. `FROM bronze_orders`), resolved against the pipeline's default catalog and schema

44. What does this Lakeflow declaration do?

```
CONSTRAINT valid_id EXPECT (id IS NOT NULL) ON VIOLATION DROP ROW
```

- A. Logs a warning for null IDs but keeps all rows
- B. Drops the entire pipeline run if any null ID is found
- C. Removes rows where `id` is `NULL` from the target and counts them in the data quality metrics
- D. Replaces `NULL` IDs with a default value

45. What happens when a Lakeflow expectation uses `ON VIOLATION FAIL UPDATE`?

- A. The failing rows are dropped and the pipeline continues
- B. The failing rows are quarantined to an error table
- C. A warning is logged but processing continues
- D. The pipeline update fails immediately, and no data is written for that table

46. What happens when an expectation is declared without an `ON VIOLATION` clause?

```
CONSTRAINT positive_amount EXPECT (amount > 0)
```

- A. Failing rows are dropped
- B. The pipeline fails
- C. All rows are kept; a warning is logged and metrics are recorded
- D. An error — `ON VIOLATION` is required

47. Which Lakeflow SQL syntax reads an upstream streaming table in the same pipeline incrementally, as a streaming source?

- A. `SELECT * FROM bronze_orders`
- B. `SELECT * FROM STREAM(bronze_orders)`
- C. `SELECT * FROM STREAMING(bronze_orders)`
- D. `SELECT * FROM READ_STREAM(bronze_orders)`

48. Where can you view data quality metrics (pass/fail counts) for Lakeflow expectations?

- A. `SHOW QUALITY METRICS FOR silver_table`
- B. `DESCRIBE EXPECTATIONS silver_table`
- C. `SELECT * FROM INFORMATION_SCHEMA.EXPECTATIONS`
- D.

```
SELECT * FROM event_log(TABLE(catalog.schema.silver_table))  
WHERE details:flow_progress:data_quality IS NOT NULL
```

49. In Lakeflow Jobs (formerly Workflows), what is a “Task” within a Job?

- A. A separate cluster configuration
- B. A single unit of work (notebook, SQL, Python script, or pipeline) that the Job orchestrates
- C. A permission setting
- D. A notification channel

50. What does `dbutils.jobs.taskValues.set()` do?

Task 1:

```
dbutils.jobs.taskValues.set(key="row_count", value=1000)
```

Task 2 (downstream):

```
count = dbutils.jobs.taskValues.get(taskKey="Task_1", key="row_count")
```

- A. Passes a value from one task to a downstream task within the same job run
- B. Sets a job-level environment variable
- C. Configures cluster parameters
- D. Sets the task retry count

51. What is a “Repair Run” in Lakeflow Jobs?

- A. A full re-run of the entire job from scratch
- B. Re-runs only the failed tasks (and their downstream dependents) without re-running successful ones
- C. A manual fix applied to the job configuration
- D. Rolling back the job to a previous version

52. Which task types can a Lakeflow Job orchestrate?

- A. Notebook, SQL, Python script, Lakeflow pipeline, JAR
- B. Only notebooks
- C. Only SQL and Python
- D. Only Lakeflow pipelines

53. Which Unity Catalog feature is attached directly to a table so that every query against that table returns only the rows the querying user is allowed to see?

- A. Column Mask
- B. Dynamic View
- C. Row Filter
- D. Data Profile

54. How do you apply a Column Mask to a table in Unity Catalog?

- A. `GRANT MASK ON COLUMN email TO mask_func`
- B. `CREATE MASK ON TABLE t COLUMN email`
- C. `UPDATE TABLE t SET COLUMN email MASKED`
- D. `ALTER TABLE t ALTER COLUMN email SET MASK catalog.schema.mask_func`

55. Which INFORMATION_SCHEMA view would you query to list all tables and views in a Unity Catalog catalog?

- A. `INFORMATION_SCHEMA.SCHEMATA`
- B. `INFORMATION_SCHEMA.TABLES`
- C. `INFORMATION_SCHEMA.COLUMNS`
- D. `INFORMATION_SCHEMA.TABLE_PRIVILEGES`

56. Which SQL function is used inside Row Filters and Column Masks to check if the querying user belongs to a specific group?

- A. `is_account_group_member('admins')`
- B. `user_in_group('admins')`
- C. `has_role('admins')`
- D. `check_group('admins')`

57. Partners upload batches of CSV files to a Unity Catalog volume at unpredictable times, a few times a day. The processing job should start soon after each delivery and should not run at all when nothing has arrived. Which trigger should the data engineer choose?

- A. A scheduled trigger that runs every 5 minutes
- B. A continuous trigger
- C. A file arrival trigger on the volume path
- D. A manual run started by the on-call engineer

58. Following least privilege, which statement grants read access on all tables in the `silver` schema to the `analysts` group?

- A. `GRANT SELECT ON SCHEMA my_catalog.silver TO analysts`
- B. `GRANT READ ON DATABASE silver TO analysts`
- C. `GRANT ALL PRIVILEGES ON SCHEMA my_catalog.silver TO analysts`
- D. `ALLOW SELECT ON my_catalog.silver FOR analysts`

59. How do you remove a Row Filter from a table?

- A. `DROP FILTER ON TABLE t`
- B. `ALTER TABLE t DROP ROW FILTER`
- C. `REVOKE ROW FILTER FROM t`
- D. `DELETE ROW FILTER ON t`

60. A data engineer works in a Git folder. The remote `main` branch is protected and only accepts reviewed pull requests. They need to change a notebook and get the change merged. What is the correct workflow?

- A. In the Git folder dialog, create a feature branch, edit the notebook, commit and push to that branch, then open a pull request in the Git provider to merge it into `main`
- B. Export the notebook as a `.dbc` archive and send it to the repository owner
- C. Edit the notebook on `main` and select **Commit & Push**; Databricks bypasses branch protection
- D. Run `databricks bundle deploy -t prod` to merge the change into `main`

61. After a colleague installs `pandas==1.5.3` as a cluster library on a shared all-purpose cluster, notebooks that worked yesterday now fail with `AttributeError` inside pandas calls. The cluster starts normally and memory metrics look healthy. What is the most likely cause and the recommended fix?

- A. Executor out-of-memory; increase `spark.executor.memory`
- B. A cluster startup failure; check the init script log
- C. Data skew; enable AQE skew join handling
- D. A library conflict with the pandas version the code expects; remove the cluster library and pin the needed version per notebook with `%pip install` (or a serverless environment)

62. A team has an external Delta table `sales.bronze.orders` registered in Unity Catalog. They want Unity Catalog to manage its storage and lifecycle (including predictive optimization) without changing the table name or losing its history. Which statement does this?

- A. `ALTER TABLE sales.bronze.orders SET TBLPROPERTIES ('managed' = 'true')`
- B. `ALTER TABLE sales.bronze.orders SET MANAGED`
- C. `CREATE TABLE sales.bronze.orders_managed DEEP CLONE sales.bronze.orders`
- D.

```
DROP TABLE sales.bronze.orders;  
CREATE TABLE sales.bronze.orders LOCATION 'abfss://landing@acct.dfs.core.windows.net/orders';
```

63. A nightly job that normally takes about 20 minutes succeeded every night this week, but today's successful run took 75 minutes. How should the data engineer confirm whether this is a regression and find where the extra time went?

- A. Restart the cluster and rerun the job; if it succeeds, no investigation is needed
- B. Run `DESCRIBE DETAIL` on the job's target table
- C. Open the job's run history and compare today's duration with previous runs, using the matrix view to see which task's duration grew compared with its usual values
- D. Read today's notebook output; durations of earlier runs are not kept

64. A data engineer must ingest data from Salesforce into Unity Catalog-governed Delta tables with minimal custom code and managed CDC handling. Which ingestion approach should they prioritize?
- A. A Lakeflow Connect managed connector
 - B. Auto Loader with `cloudFiles` on exported CSV files
 - C. A JDBC driver in a notebook scheduled with Lakeflow Jobs
 - D. COPY INTO from a cloud storage staging area
65. A team has one codebase for a Lakeflow job and pipeline that must be promoted across dev, test and prod workspaces with different catalogs. Which mechanism is designed for this?
- A. Cloning the Git folder into each workspace and editing notebook paths
 - B. A single job with one task per environment
 - C. Exporting the job JSON and importing it manually in each workspace
 - D. Declarative Automation Bundle (formerly Databricks Asset Bundles) targets with variable overrides
66. In the Spark UI, one stage shows most tasks finishing in seconds while a few tasks run for many minutes and shuffle read for those tasks is far larger than the median. What is the most likely cause?
- A. Disk spill caused by too much executor memory
 - B. A library conflict on the cluster
 - C. Data skew — a few partitions hold disproportionately many records for certain keys
 - D. Too few columns selected in the query
67. An organization wants column masking applied centrally to EVERY table column tagged `pii = email`, without adding a mask to each table one by one. Which Unity Catalog feature does this?
- A. ABAC policies on governed tags
 - B. `DENY` privileges on the affected columns
 - C. Dynamic views with `is_account_group_member()`
 - D. Delta Sharing recipient restrictions
68. According to current Databricks guidance, what is the recommended default way to keep a Unity Catalog managed Delta table well clustered and compacted over time?
- A. A nightly job running `OPTIMIZE ... ZORDER BY` and weekly `VACUUM`
 - B. Manual repartitioning with `df.repartition()` before every write
 - C. Hive-style static partitioning on a date column
 - D. Predictive optimization with (Automatic) Liquid Clustering, e.g. `CLUSTER BY AUTO`

Answer key

Day	Answers
Day 1	1 C, 2 A, 3 C, 4 D, 5 C, 6 A, 7 B, 8 A, 9 B, 10 A
Day 1	11 D, 12 C, 13 D, 14 B, 15 D, 16 B, 17 D, 18 A, 19 C, 20 B
Day 2	21 A, 22 C, 23 B, 24 C, 25 A, 26 B, 27 C, 28 B, 29 D, 30 C
Day 2	31 A, 32 D, 33 B, 34 D, 35 A, 36 B, 37 D, 38 C, 39 A, 40 D
Day 3	41 C, 42 B, 43 D, 44 C, 45 D, 46 C, 47 B, 48 D, 49 B, 50 A
Day 3	51 B, 52 A, 53 C, 54 D, 55 B, 56 A, 57 C, 58 A, 59 B, 60 A
Day 3	61 D, 62 B, 63 C, 64 A, 65 D, 66 C, 67 A, 68 D

The questions reflect the platform as of September 2026 and the Databricks Certified Data Engineer Associate exam guide of May 2026. Product names follow current Databricks terminology: Lakeflow Jobs (formerly Workflows), Lakeflow Spark Declarative Pipelines (formerly Delta Live Tables) and Declarative Automation Bundles (formerly Databricks Asset Bundles). The same questions, with explanations for every answer, are in the course notebooks `utilization/quiz/day1_quiz.ipynb`, `day2_quiz.ipynb` and `day3_quiz.ipynb`.